

Simulations of Ballot Polling Risk-Limiting Audits

No Author Given

No Institute Given

Abstract. In this paper we present simulation results comparing the risk, stopping probability and number of ballots required over multiple rounds of ballot-polling risk limiting audits (RLAs) MINERVA, Selection-Ordered BRAVO and End-of-Round BRAVO. We also present details on the R2B2 open source software library and the related simulation software.

BRAVO is the most commonly used ballot-polling RLA. MINERVA was recently proposed, and requires fewer first-round ballots, on average, than both Selection-Ordered BRAVO and End-of-Round BRAVO when the first-round stopping probability is large.

An open question, however, is how MINERVA compares to Selection-Ordered BRAVO and End-of-Round BRAVO over multiple rounds. In this paper, we present results from simulations of multiple round audits with first-round stopping probabilities of 0.9, a common choice among election officials. Because the size of rounds in MINERVA needs to be predetermined (and independent of the audit draws), we use two pre-determined round sequences for MINERVA: (a) all rounds are of the same size (which might be a practical choice for election officials if they had planned to have resources for the first round size) and (b) each round is 1.5 times the previous one (which is the preset value for MINERVA as integrated into election audit software Arlo).

We show that the simulation results are consistent with predictions of the R2B2 open-source library for ballot polling audits. We also observe that both BRAVO audits are more conservative than needed, while MINERVA audits stop with fewer ballots.

Keywords: Risk-limiting audit · Ballot polling audit

1 Introduction

Need motivating paragraph and two sentence summary of results here.

1.1 Background

In the general ballot-polling risk limiting audit (RLA), a number of ballots are drawn and tallied in what is termed a *round* of ballots [5]. A statistical measure is then computed to determine whether there is sufficient evidence to declare the election outcome correct within the pre-determined risk limit. Because the

decision is made after drawing a round of ballots, the audit is termed a *round-by-round (R2)* audit. The special case when round size is one—that is, stopping decisions are made after each ballot draw—is a *ballot-by-ballot (B2)* audit.

The BRAVO audit [1] is designed for use as a B2 audit: it requires the smallest expected number of ballots when the true tally of the underlying election is as announced, and stopping decisions are made after each ballot draw. In practice, election officials draw many ballots at once, and the BRAVO stopping rule needs to be modified for use in an R2 audit that is not B2. There are two obvious approaches. The B2 stopping condition can be applied once at the end of each round: End-of Round (EoR) BRAVO. Alternatively, the order of ballots in the sample can be tracked by election officials and the B2 BRAVO stopping condition can be applied retroactively after each ballot drawn: Selection-Ordered (SO) BRAVO. SO BRAVO requires fewer ballots on average than EoR BRAVO but requires the work of tracking the order of ballots rather than just their tally.

MINERVA was designed for R2 audits and applies its stopping rule once for each round. Thus it does not require the tracking of ballots that SO BRAVO does. It has been proven that MINERVA is a risk-limiting audit and requires fewer ballots to be sampled than EoR BRAVO when an audit is performed in rounds. Computations of first-round size for a 0.9 stopping probability when the election is as announced have been computed for a wide range of margins and shown to be smaller than those for both EoR and SO BRAVO. First round simulations of MINERVA [4] demonstrate that its first-round properties—regarding the probabilities of stopping when the underlying election is tied and when it is as announced—are as predicted for first round sizes with stopping probability 0.9. There are no results, either theoretical or based on simulations, regarding the expected number of ballots drawn in a MINERVA audit. Further, there is no literature comparing simulations of MINERVA and EoR or SO BRAVO, or studying the sizes of multiple-round MINERVA audits.

Both BRAVO and MINERVA have been integrated into election audit software *Arlo* [3].

1.2 Our Results

1.3 Organization

2 Related work

3 Definitions

Now we present relevant definitions. We consider a two candidate plurality contest. To begin we define an audit. We assume that ballots are sampled with replacement.

Definition 1. *An audit \mathcal{A} takes a sample X as input and gives one of the following decisions*

1. *Correct: the audit is complete*

2. *Uncertain: continue the audit*

All of the audits discussed in this paper are modeled as binary hypothesis tests. Under the alternate hypothesis, H_a , the announced outcome is correct. That is, the true underlying ballot distribution is given by the announced ballot tallies. Under the null hypothesis, H_0 , the true outcome is a tie (or the announced winner lost by one vote, and the number of ballots is large enough that the probability of drawing a ballot for the winner is that of drawing one for the winner). Now we define key attributes of an audit that we will consider while analyzing our simulations. The maximum risk of an audit is the probability that an audit stops, given that the underlying election is a tie.

Definition 2 (Risk). *The maximum risk R of an audit \mathcal{A} is*

$$R(\mathcal{A}) = \Pr[\mathcal{A}(X) = \text{Correct} \mid H_0]$$

This leads us to the following simple definition of an α -RLA.

Definition 3 (Risk Limiting Audit (α -RLA)). *An audit \mathcal{A} is a Risk Limiting Audit with risk limit α iff*

$$R(\mathcal{A}) \leq \alpha.$$

It is useful to discuss the probability of an audit stopping in the j^{th} round, given that the underlying election is as announced.

Definition 4 (Stopping Probability). *The stopping probability S of an audit \mathcal{A} in round j is*

$$S_j(\mathcal{A}) = \Pr[\mathcal{A}(X) = \text{Correct in round } j \wedge \mathcal{A}(X) \neq \text{Correct previously} \mid H_a]$$

The notion of stopping probability can be useful for selecting round sizes.

4 Software

In this section we describe the software implementing ballot polling audits, termed the R2B2 library, and the simulator software used for this research. All the software is released as open source under the MIT License.

4.1 R2B2 Library

The R2B2 Python library [2] provides a framework for the exploration of round-by-round and ballot-by-ballot RLAs. The goal in designing R2B2 is two fold:

1. Provide an elegant Python library which can be easily imported and used in any other code base.
2. Provide an interactive set of tools which can be utilized ‘out-of-the-box’ for experimenting with and learning about risk-limiting audits.

Design The high-level design of R2B2 is an object-oriented view of election audits. The three main object classes, **Election**, **Contest**, and **Audit**, serve to group data into logically independent structures.

The **Election** contains the information that comprises an entire election, most importantly, the total number of ballots cast in the election and the list of **Contests** from the election. At the moment **Election** does not offer functionality beyond grouping **Contests**.

The **Contest** contains the information related to a single contest such as the ballots cast in that contest, the candidates, the type of contest, and the reported tally. Providing a structure to hold this information independent of any particular audit is especially useful for exploratory work.

The **Audit** contains information related to the audit parameters for a single contest, such as the risk limit, sampling method, and **Contest** to audit. It is important to note the **Audit** is an Abstract Base Class upon which specific RLAs are built. It only contains the parameters and attributes common to the RLAs of this paper and provides a set of methods that can be called by any audit implementation. The functionality of **Audit** can be divided into two basic groups: *interactive* and *bulk*.

The interactive implementation allows users to execute an audit step-by-step as it might progress during a live election audit through the following:

- The `run()` method begins an interactive audit executing where users are prompted for round sizes and the counts of winner ballots found in the sample and in return are given information about the current state of the audit and whether the stopping condition(s) have been met.
- Two distributions representing the null and alternative hypotheses are maintained and allow for computation of the audits per-round risk and stopping probability schedules.
- Before each round, the audit will recommend possible next round sizes given different criteria, such as a set of desired stopping probabilities.

The bulk implementations allows users to generate a larger set of data from an audit such as:

- A set of stopping conditions given a set of round sizes.
- A set of risk levels given a set of round size and winner ballots pairs.
- A list of all stopping conditions from the minimum to the maximum round size.

Usage R2B2 makes understanding and exploring election audits simple for the user with no Python knowledge while simultaneously providing a comprehensive set of tools for the experienced Python developer.

Using R2B2 is as simple as using any other Python library: simply import the library and all of the functionality is at your finger tips. Not only does this allow users to write their own Python scripts for exploring RLAs, it also allows

R2B2 to be plugged in to any other Python library. See the following Jupyter Notebooks for information on the usage of R2B2:

- Basic Usage
- Generating Graphs

R2B2 also provides a significant amount of functionality ‘out-of-the-box’ for educational or exploratory use. For those who wish learn about RLAs without having to write any code themselves, R2B2 provides a command line tool for both interactive auditing and generating audit results and statistics for larger data sets.

4.2 Simulation Software

As described above, our open source audit software library R2B2 has implementations of several ballot polling risk-limiting audits as well as a simulator, all written in Python. For each of these audits, the software can compute the stopping condition for a given sample and estimates of the next round size to achieve a desired stopping probability. For a given audit and random seed, the simulator draws random samples, with replacement, using a pseudorandom number generator, [need to check]. given the number of votes for each candidate, and the number of invalid votes, in the underlying election (these need not be chosen to be those announced).

When the number of candidates is more than two, the audit is carried out pairwise for each candidate pair, and votes for all other candidates are considered invalid votes. It is often useful to consider two underlying elections: one corresponding to a tie between the pair, and the other corresponding to the results as announced. These are the probability distributions represented by the null and alternative hypotheses.

After drawing a simulated sample of ballots, the simulator evaluates the given audit’s stopping condition for this sample. If the audit stops, the simulation stops, and if the audit continues, the simulation draws another round. The abstract simulator class does not prescribe any one method for choosing round sizes. We implement several classes to support various round size choices: round sizes from an estimate to achieve a desired probability of stopping, predetermined round sizes, and pseudorandomly-generated round sizes.

4.3 Testing

The R2B2 (python) software is used to compute stopping conditions and next round estimates, it is intended for use by us and other researchers, and designed for this purpose. We have also independently implemented all the functionality in matlab (the two codebases are written by different individuals) and have extensively checked the results of the two codebases. Additionally, for use in regular election audits by election officials, we have written an add-on to the *Arlo* risk-limiting audit software, which has also been extensively checked against these two codebases.

5 Simulation Results

For this paper, we simulated audits for the 2020 Presidential election in all US states whose margin was at least 5%. Round sizes increase roughly proportional to the inverse square of the margin, so smaller margins are computationally much more expensive to simulate. For each of these states, we simulated $10,000 = 10^4$ audits with the announced underlying ballot distribution and an additional $10,000 = 10^4$ audits with a tie as the underlying ballot distribution. These are reasonable choices for initial simulation experiments because most audits frame the stopping decision as a binary hypothesis test where the null hypothesis assumes an underlying tie and the alternative hypothesis assumes the announced distribution. A standard first round size in ballot-polling audits has been one which achieves a 90% probability of stopping, and we chose round sizes to reflect this standard in our simulations. We ran our simulations for up to five rounds. With a 90% stopping probability, an audit only has a α probability of not stopping by the end of the fifth round, assuming the outcome was correctly announced.

5.1 End-of-Round Bravo

For the EoR BRAVO simulations, our software estimated and used for each round the minimum round size that would achieve a 90% probability of stopping. For the simulations with the announced outcome as the underlying ballot distribution, for each round j we computed the proportion of audits that stopped in round j among those that had not stopped before round j .




Fig. 1. For the first three rounds of the EoR BRAVO simulations, this plot shows the proportion of audits that stopped in the j th round to all audits which had not yet stopped before the j th round.

In Figure 1, we display proportions for only the first three rounds since very few audits, $(.1)^{j-1} \cdot (10^4)$ on average, make it to the j th round. As a result, the proportions are based on an exponentially smaller dataset in each round. The proportions shown in Figure 1 give an estimate of the true probability of an EOR BRAVO audit stopping in the j th round, given that it has not already stopped in a previous round. In 1, we see that, especially in earlier rounds for which the values are more representative of true audit behavior, our predictions are accurate. In particular, the average across all margins is just above $.9 = 90\%$ for all three rounds.

The risk of this audit, across all 5 rounds, is an important metric since it determines whether an audit is risk-limiting. Therefore, we now consider the proportion of audits that stopped with an underlying tie. This proportion, for a risk-limiting audit, should approach a value less than the risk limit, 0.1 , as more audits are performed.

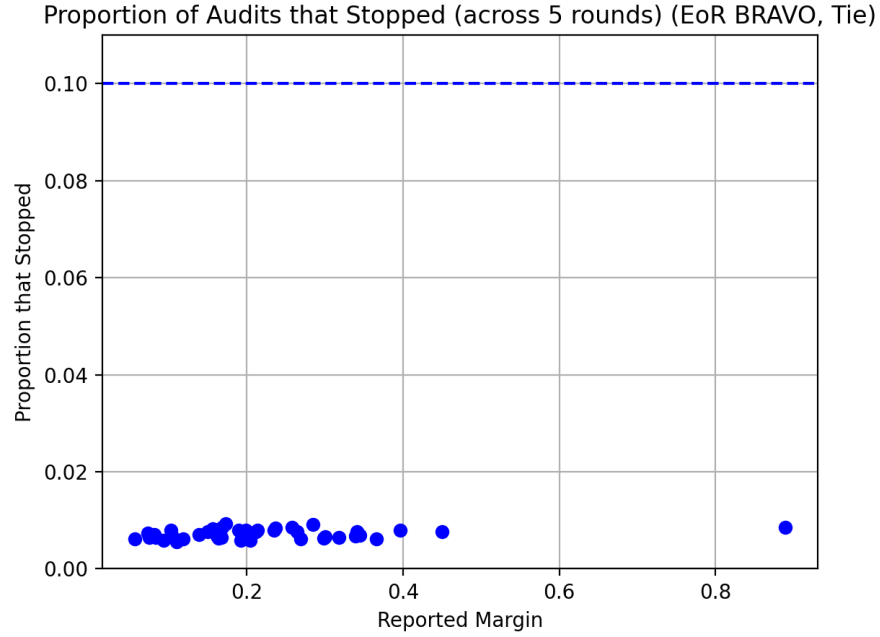


Fig. 2. For each state margin, this plot shows the proportion of EOR BRAVO audits with an underlying tie that stopped.

Figure 2 makes it clear that the risk of EOR BRAVO is roughly an order of magnitude less than the risk limit. That is, on average, roughly 10 times more audits could have stopped with the audit still meeting the risk limit. These simulations support the claim made in [MINERVA paper] that EOR BRAVO is unnecessarily conservative and thus requires more ballots on average.

5.2 Selection-Ordered Bravo

For the SO BRAVO simulations, our software estimated and used for each round the minimum round size that would achieve a 90% probability of stopping. For the simulations with the announced outcome as the underlying ballot distribution, for each round j we computed the proportion of audits that stopped in round j among those that had not stopped before round j .




Fig. 3. For the first three rounds of the SO BRAVO simulations, this plot shows the proportion of audits that stopped in the j th round to all audits which had not yet stopped before the j th round.

In Figure 3, we again display proportions for only the first three rounds. The proportions shown in Figure 3, like those in 1, give an estimate of the true probability of an SO BRAVO audit stopping in the j th round, given that it has not already stopped in a previous round. In 3, we see that our round size predictions are relatively accurate, all three rounds being near 90%.

Again, we will next consider the proportion of audits that stopped with an underlying tie over all five rounds. This proportion, for a risk-limiting audit, should approach a value less than the risk limit, 0.1, as more audits are performed.




Fig. 4. For each state margin, this plot shows the proportion of SO BRAVO audits with an underlying tie that stopped in any of the five rounds.

In Figure 4 we show only the results for the 13 states whose simulations with an underlying tie have finished running. To estimate the minimum next round size that achieves a desired stopping probability, the SO BRAVO software generates the probability distribution ballot by ballot since the stopping condition needs to be evaluated for each individual ballot drawn. This means, when an underlying tie causes audits to quickly move on to larger rounds, the simulations are computationally expensive. Regardless, the claim that these simulations support, that SO BRAVO is a Risk-Limiting Audit, is a proven and well known result. Notice in Figure 4, that the risk of SO BRAVO is much nearer the risk limit than that of EoR BRAVO.

5.3 Minerva Simulations

For MINERVA, it has not been shown that round sizes can be chosen during the audit. That is, an adversary with knowledge of the history of the audit may be able to choose round sizes which cause the risk of the audit to exceed the risk limit. For this reason, we have to choose the round sizes of a MINERVA audit a priori. For this paper, we consider two choices of round sizes. For both, we estimate and then use the minimum first round size which achieves a 90%

probability of stopping. Then, for subsequent rounds, we either (i) draw the same number of ballots in each round or (ii) multiply the previous round size by a factor of 1.5 and sample this many new ballots. We consider the case of drawing samples of the same size because it may reflect a practical way to continue an audit; if election officials have selected some first round size within reasonable logistical bounds, drawing the same number of ballots in subsequent rounds may be practical. We also consider round sizes with samples increasing by a multiple of 1.5 because this multiple gives a very rough approximation of round sizes with a 90% probability of stopping.

Round Sizes with Multiple of 1.0 As with the preceding simulations, we ran $10,000 = 10^4$ trials per state for both the underlying tie and underlying reported outcome.




Fig. 5. This plot shows, for each state’s margin, the proportion of audits that stopped of all 10^4 audits with an underlying tie in the simulations with a round size multiple of 1.0.

Notice in Figure 5 that all simulations for a tie had a proportion of audits that stopped less than .1, the risk limit, supporting the claim that MINERVA is a risk limiting audit. Unlike EOR BRAVO, the experimental risks here are

much closer to the risk limit, showing that MINERVA stops on average with a less conservative risk; MINERVA is sharper.




Fig. 6. This plot shows, for each state’s margin, the proportion of audits that stopped of all 10^4 audits with the announced results as the underlying distribution in the MINERVA simulations with a round size multiple of 1.0.

Figure 6 shows that the round size estimate for the first round achieved the desired stopping probability of 90%. For subsequent rounds, the multiplier of 1.5 did not consistently achieve 90% stopping probability, but it was within roughly 10%.




Fig. 7. This plot shows, for each state's margin, the proportion of audits that stopped of all 10^4 audits with an underlying tie in the simulations with a round size multiple of 1.5.

Round Sizes with Multiple of 1.5 Figure 7 shows, for each state's margin, the proportion of audits that stopped of all 10^4 audits with an underlying tie in the simulations with a round size multiple of 1.5.




Fig. 8. This plot shows, for each state’s margin, the proportion of audits that stopped in each round with the announced results as the underlying distribution in the MINERVA simulations with a round size multiple of 1.5.

Figure 8 shows that the round size estimate for the first round achieved the desired stopping probability of 90%. For subsequent rounds, the multiplier of 1.5 did not consistently achieve 90% stopping probability, but it was within roughly 10%.

6 Conclusions and Future Work

References

1. Lindeman, M., Stark, P.B., Yates, V.S.: BRAVO: Ballot-polling risk-limiting audits to verify outcomes. In: EVT/WOTE (2012)
2. Morin, S., McClearn, G.: The R2B2 (Round-by-Round, Ballot-by-Ballot) library, <https://github.com/gwexploratoryaudits/r2b2>
3. VotingWorks: Arlo, <https://voting.works/risk-limiting-audits/>
4. Zagórski, F., McClearn, G., Morin, S., McBurnett, N., Vora, P.L.: The athena class of risk-limiting ballot polling audits. CoRR **abs/2008.02315** (2020), <https://arxiv.org/abs/2008.02315>

5. Zagórski, F., McClearn, G., Morin, S., McBurnett, N., Vora, P.L.: Minerva— an efficient risk-limiting ballot polling audit. In: 30th USENIX Security Symposium (USENIX Security 21). pp. 3059–3076. USENIX Association (Aug 2021), <https://www.usenix.org/conference/usenixsecurity21/presentation/zagorski>